

## SUMMARY

### DataFrame dfa -

Data with dummy variables for Location, Category, Worktype and Job Class (Data Scientist, Data Analyst, Business Intelligence). Training dataset was created with rows where FinalSal is NOT NULL. Test dataset was created with rows where FinalSal IS NULL. MLR OLS regression was run for three scenarios

1. Features = Location, Job Class, Category, Work Type MAE 44301, MSA 6014002312, R2 0.044
2. Features = Location, Job Class, Work Type MAE 40480, MSA 5322109879, R2 0.154 >> best and overall best
3. Features = Location, Job Class MAE 45156, MSA 6220574795, R2 0.011 Visual inspection of the plots also indicate that Scenario 2 is the best.

### DataFrame dfn -

Data with numeric values mapped for all classifier datapoints. No dummy variables created. Training dataset was created with rows where FinalSal is NOT NULL. Test dataset was created with rows where FinalSal IS NULL.

KNN regression was run for three scenarios Accuracy scores were taken for Odd number of neighbours between 1 and 51.

1. Features = Location, Job Class, Category, Work Type MAE 44152, MSA 6005511961, R2 0.045 Neighbors 21 >> best
2. Features = Location, Job Class, Work Type MAE 44568, MSA 6223339862, R2 0.010 Neighbors 27
3. Features = Location, Job Class MAE 44568, MSA 6223339862. R2 0.011 Neighbors 27

CART regression was run for three scenarios Accuracy scores were taken for Depth between 1 and 7 and None.

1. Features = Location, Job Class, Category, Work Type MAE 43354, MSA 5768771897, R2 0.083 Depth 2 >> best
2. Features = Location, Job Class, Work Type MAE 44867, MSA 6159154547, R2 0.020 Depth 4
3. Features = Location, Job Class MAE 44867, MSA 6159154547, R2 0.021 Depth 4

Enhancements if there is time

1. Use bootstrapping to increase size of Train database and check accuracy of OLS model

In [4]:

```
import requests
from bs4 import BeautifulSoup as bs
import pandas as pd
import numpy as np
import csv
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

## LOAD DATA

In [5]:

```
# This is data where dummy variables have been created.
```

```
dfe = pd.read_csv('./Data/dfe.csv')  
dfe.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1286 entries, 0 to 1285  
Data columns (total 21 columns):  
Unnamed: 0      1286 non-null int64  
jobid           1286 non-null int64  
finalsal        279 non-null float64  
category_HR      1286 non-null int64  
category_IT      1286 non-null int64  
category_MC      1286 non-null int64  
category_SAL     1286 non-null int64  
category_ST      1286 non-null int64  
location_ADE     1286 non-null int64  
location_BRI     1286 non-null int64  
location_HOB     1286 non-null int64  
location_MEL     1286 non-null int64  
location_PER     1286 non-null int64  
location_SYD     1286 non-null int64  
jobclass_BI      1286 non-null int64  
jobclass_DA      1286 non-null int64  
jobclass_DS      1286 non-null int64  
worktype_CAS     1286 non-null int64  
worktype_CON     1286 non-null int64  
worktype_FT      1286 non-null int64  
worktype_PT      1286 non-null int64  
dtypes: float64(1), int64(20)  
memory usage: 211.1 KB
```

In [6]:

```
# This is data where salaries are not null and dummy variables have been created.
train = dfe[dfe.finalsal.notnull()]
# This is data where salaries are null and dummy variables have been created.
test = dfe[dfe.finalsal.isnull()]
test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1007 entries, 1 to 1284
Data columns (total 21 columns):
Unnamed: 0      1007 non-null int64
jobid          1007 non-null int64
finalsal       0 non-null float64
category_HR    1007 non-null int64
category_IT    1007 non-null int64
category_MC    1007 non-null int64
category_SAL   1007 non-null int64
category_ST    1007 non-null int64
location_ADE   1007 non-null int64
location_BRI   1007 non-null int64
location_HOB   1007 non-null int64
location_MEL   1007 non-null int64
location_PER   1007 non-null int64
location_SYD   1007 non-null int64
jobclass_BI    1007 non-null int64
jobclass_DA    1007 non-null int64
jobclass_DS    1007 non-null int64
worktype_CAS   1007 non-null int64
worktype_CON   1007 non-null int64
worktype_FT    1007 non-null int64
worktype_PT    1007 non-null int64
dtypes: float64(1), int64(20)
memory usage: 173.1 KB
```

## Multivariate Linear Regression

In [8]:

```
import statsmodels.api as sm
```

C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\statsmodels\compat\pandas.py:56: FutureWarning: The pandas.core.datetools module is deprecated and will be removed in a future version. Please use the pandas.tseries module instead.

```
from pandas.core import datetools
```

In [9]:

```
# Setup features for modelling with Ordinary Least Squares
X = train[['category_IT', 'category_MC', 'category_SAL', 'category_ST', \
          'location_BRI', 'location_HOB', 'location_MEL', 'location_PER', 'location_SYD', \
          'jobclass_BI', 'jobclass_DA', 'jobclass_DS', 'worktype_CON', 'worktype_FT', 'worktype_
Y = train['finalsal']
mlrmod1 = sm.OLS(Y,X).fit()
y_pred = mlrmod1.predict(X)
```

In [123]:

```
print('MAE', mean_absolute_error(Y, y_pred))
print('MSA', mean_squared_error(Y, y_pred))
print('R2', r2_score(Y, y_pred))
```

```
MAE 44300.8903268761
MSA 6014002312.054384
R2 0.04403437951642675
```

In [11]:

```
# Get errors / residuals and predicted into dataframe
train['residual'] = Y - y_pred
train['pred'] = y_pred
```

```
C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\ipykernel_launcher.py:1:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

"""Entry point for launching an IPython kernel.

```
C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\ipykernel_launcher.py:2:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```


See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

In [12]:

```
# Plot actual salary and predicted salary  
sns.lmplot(x='finalsal', y='pred', data=train)
```

Out[12]:

<seaborn.axisgrid.FacetGrid at 0x2b501afb588>




In [13]:

```
# Plot actual salary and residuals  
sns.lmplot(x='finalsal', y='residual', data=train)
```

Out[13]:

<seaborn.axisgrid.FacetGrid at 0x2b501c07588>



In [14]:

```
# Remove category from features
X1 = train[['location_BRI', 'location_HOB', 'location_MEL', 'location_PER', 'location_SYD', \
            'jobclass_BI', 'jobclass_DA', 'jobclass_DS', 'worktype_CON', 'worktype_FT', 'worktype_
Y1 = train['finalsal']
mlrmod2 = sm.OLS(Y1, X1).fit()
y_pred1 = mlrmod2.predict(X1)
```

In [15]:

```
print('MAE', mean_absolute_error(Y1, y_pred1))
print('MSA', mean_squared_error(Y1, y_pred1))
print('R2', r2_score(Y1, y_pred1))
```

```
MAE 40480.24195767891
MSA 5322109878.8477125
R2 0.15401527824214778
```

In [17]:

```
train['residual1'] = Y1 - y_pred1
train['pred1'] = y_pred1
```

C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\ipykernel\_launcher.py:3:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

This is separate from the ipykernel package so we can avoid doing imports until

C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\ipykernel\_launcher.py:4:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)


after removing the cwd from sys.path.

In [19]:

```
# Linearity plot  
sns.lmplot(x='finalsal', y='pred1', data=train)
```

Out[19]:

<seaborn.axisgrid.FacetGrid at 0x2b502d1fb38>




In [20]:

```
# Equality of variance.  
sns.lmplot(x='finalsal', y='residual1', data=train)
```

Out[20]:

<seaborn.axisgrid.FacetGrid at 0x2b502d851d0>




In [21]:

```
# Normality of errors
sns.distplot(train['residual1'])
```

Out[21]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x2b502dc00b8>



In [22]:

```
# Remove worktype from features
X2 = train[['location_BRI', 'location_HOB', 'location_MEL', 'location_PER', 'location_SYD', \
            'jobclass_BI', 'jobclass_DA', 'jobclass_DS']]
Y2 = train['finalsal']
mlrmod3 = sm.OLS(Y2, X2).fit()
y_pred2 = mlrmod3.predict(X2)
```

In [23]:

```
print('MAE', mean_absolute_error(Y2, y_pred2))
print('MSA', mean_squared_error(Y2, y_pred2))
print('R2', r2_score(Y2, y_pred2))
```

```
MAE 45155.65857510886
MSA 6220574795.513222
R2 0.01119831094213597
```

## KNN Classification

In [24]:

```
# Load data with numeric values mapped for all classifier datapoints. No dummy variables created
dfn = pd.read_csv('./Data/dfn.csv')
```



In [25]:

```
dfn.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1286 entries, 0 to 1285
Data columns (total 7 columns):
Unnamed: 0      1286 non-null int64
jobid          1286 non-null int64
category       1286 non-null int64
location       1286 non-null int64
jobclass       1286 non-null int64
finalsal       279 non-null float64
worktype       1286 non-null int64
dtypes: float64(1), int64(6)
memory usage: 70.4 KB
```

In [26]:

```
ktrain = dfn[dfn.finalsal.notnull()]
ktest = dfn[dfn.finalsal.isnull()]
```

In [27]:

```
# Loading Library
from sklearn.cross_validation import train_test_split, cross_val_score
from sklearn.neighbors import KNeighborsClassifier, KNeighborsRegressor
from sklearn.metrics import accuracy_score
```

C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\sklearn\cross\_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model\_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.  
"This module will be removed in 0.20.", DeprecationWarning)

In [30]:

```
kX = ktrain[['category', 'location', 'jobclass', 'worktype']]
kY = ktrain['finalsal']
ktX = ktest[['category', 'location', 'jobclass']]
ktY = ktest['finalsal']
```

In [31]:

```
# Check for model accuracy
sc=[]
for k in range(1,51,2):
    x=[]
    knn = KNeighborsRegressor(n_neighbors=k)
    knn.fit(kX,kY)
    x=[k,knn.score(kX,kY)]
    sc.append(x)
```

In [43]:

```
# Neighbours value of 21 gives the best accuracy
knn = KNeighborsRegressor(n_neighbors=21)
knn.fit(kX,kY)
predict = knn.predict(kX)
```

In [45]:

```
print('MAE', mean_absolute_error(kY, predict))
print('MSA', mean_squared_error(kY, predict))
print('R2', r2_score(kY, predict))
```

```
MAE 44152.29527223075
MSA 6005511961.117378
R2 0.04538397720203946
```

In [58]:

```
kX1 = ktrain[['location','jobclass','worktype']]
kY1 = ktrain['finalsal']
sc=[]
for k in range(1,51,2):
    x=[]
    knn1 = KNeighborsRegressor(n_neighbors=k)
    knn1.fit(kX1,kY1)
    x=[k,knn1.score(kX1,kY1)]
    sc.append(x)
#sc
```

In [59]:

```
knn1 = KNeighborsRegressor(n_neighbors=27)
knn1.fit(kX1,kY1)
predict1 = knn1.predict(kX1)
```

In [60]:

```
print('MAE', mean_absolute_error(kY1, predict1))
print('MSA', mean_squared_error(kY1, predict1))
print('R2', r2_score(kY1, predict1))
```

```
MAE 44568.549449090664
MSA 6223339861.900076
R2 0.010758785270706173
```

In [62]:

```
kX2 = ktrain[['location','jobclass']]
kY2 = ktrain['finalsal']
sc=[]
for k in range(1,51,2):
    x=[]
    knn2 = KNeighborsRegressor(n_neighbors=k)
    knn2.fit(kX2,kY2)
    x=[k,knn2.score(kX2,kY2)]
    sc.append(x)
#sc
```

In [63]:

```
# [27, 0.010758785270706173],  
knn2 = KNeighborsRegressor(n_neighbors=27)  
knn2.fit(kX2,kY2)  
predict2 = knn2.predict(kX2)
```

In [64]:

```
print('MAE', mean_absolute_error(kY2, predict2))  
print('MSA', mean_squared_error(kY2, predict2))  
print('R2', r2_score(kY2, predict2))
```

```
MAE 44568.549449090664  
MSA 6223339861.900076  
R2 0.010758785270706173
```

In [104]:

```
ktrain['predict'] = predict  
ktrain['knndiff'] = ktrain['finalsal'] - ktrain['predict']
```

```
C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\ipykernel_launcher.py:1:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

"""Entry point for launching an IPython kernel.

```
C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\ipykernel_launcher.py:2:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```


See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

In [102]:

```
sns.lmplot(x='finalsal', y='predict', data=ktrain)
```

Out[102]:

<seaborn.axisgrid.FacetGrid at 0x1d9b838d518>




In [105]:

```
sns.lmplot(x='finalsal', y='knndiff', data=ktrain)
```

Out[105]:

<seaborn.axisgrid.FacetGrid at 0x1d9b8e494e0>



## CART Regression Decision Tree

In [66]:

```
from sklearn.tree import DecisionTreeRegressor
```

In [67]:

```
# kX and kY variables
# Get the best depth
sc=[]
for d in [1,2,3,4,5,6,7,None]:
    dtr = DecisionTreeRegressor(max_depth=d)
    dtr.fit(kX, kY)
    dtr_scores = cross_val_score(dtr, kX, kY, cv=3)
    print(d, np.mean(dtr_scores))
```

```
1 -0.041955806818952356
2 -0.02580911105425221
3 -0.037132341316682194
4 -0.07000370264515099
5 -0.05102493464926883
6 -0.051050275369455576
7 -0.05930228972147781
None -0.0707021316633425
```

In [72]:

```
# Depth 2 is the best score
dtr = DecisionTreeRegressor(max_depth=4)
dtr.fit(kX, kY)
predictcart = dtr.predict(kX)
```

In [73]:

```
print('MAE', mean_absolute_error(kY, predictcart))
print('MSA', mean_squared_error(kY, predictcart))
print('R2', r2_score(kY, predictcart))
```

```
MAE 43353.78962406638
MSA 5768771897.200209
R2 0.08301538310327672
```

In [78]:

```
sc=[]
for d in [1,2,3,4,5,6,7,None]:
    dtr1 = DecisionTreeRegressor(max_depth=d)
    dtr1.fit(kX1, kY1)
    dtr_scores = cross_val_score(dtr1, kX1, kY1, cv=3)
    print(d, np.mean(dtr_scores))
```

```
1 -0.028327693066799275
2 -0.02975594655907791
3 -0.05755715832111874
4 -0.057774538046887626
5 -0.057774538046887626
6 -0.057774538046887626
7 -0.057774538046887626
None -0.057774538046887626
```

In [80]:

```
# Depth 4 is the best score
dtr1 = DecisionTreeRegressor(max_depth=4)
dtr1.fit(kX1, kY1)
predictcart1 = dtr1.predict(kX1)
```

In [81]:

```
print('MAE', mean_absolute_error(kY1, predictcart1))
print('MSA', mean_squared_error(kY1, predictcart1))
print('R2', r2_score(kY1, predictcart1))
```

```
MAE 44866.632749691715
MSA 6159154547.129303
R2 0.020961467457540017
```

In [82]:

```
sc=[]
for d in [1,2,3,4,5,6,7,None]:
    dtr2 = DecisionTreeRegressor(max_depth=d)
    dtr2.fit(kX2, kY2)
    dtr_scores = cross_val_score(dtr2, kX2, kY2, cv=3)
    print(d, np.mean(dtr_scores))
```

```
1 -0.028327693066799275
2 -0.02975594655907791
3 -0.05755715832111874
4 -0.057774538046887626
5 -0.057774538046887626
6 -0.057774538046887626
7 -0.057774538046887626
None -0.057774538046887626
```

In [83]:

```
# Depth 4 is the best score
dtr2 = DecisionTreeRegressor(max_depth=4)
dtr2.fit(kX2, kY2)
predictcart2 = dtr2.predict(kX2)
```

In [84]:

```
print('MAE', mean_absolute_error(kY2, predictcart2))
print('MSA', mean_squared_error(kY2, predictcart2))
print('R2', r2_score(kY2, predictcart2))
```

```
MAE 44866.632749691715
MSA 6159154547.129303
R2 0.020961467457540017
```

In [106]:

```
ktrain['predictcart'] = predictcart
ktrain['cartdiff'] = ktrain['finalsal'] - ktrain['predictcart']
```

C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\ipykernel\_launcher.py:1:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

"""Entry point for launching an IPython kernel.

C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\ipykernel\_launcher.py:2:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead


See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

In [103]:

```
sns.lmplot(x='finalsal', y='predictcart', data=ktrain)
```

Out[103]:

<seaborn.axisgrid.FacetGrid at 0x1d9b8dddef0>




In [107]:

```
sns.lmplot(x='finalsal', y='cartdiff', data=ktrain)
```

Out[107]:

<seaborn.axisgrid.FacetGrid at 0x1d9b8e96390>



In [ ]: