

## Import modules ¶

### SUMMARY

Data was scraped from seek.com.au and the following datapoints were extracted in relation to job posted on the website: category Job Category company Company placing the Ad jobclass Type of job - Data Science, Data Analyst and Business Intelligence jobdate Data Job posted jobid Job Id jobtext Free text on the jobsite location Job location eg: Melbourne subcategory Job Subcategory suburb Job suburb eg: Melbourne CBD title Job Title worktype Work Type eg: Full Time

The following tasks are performed:

1. Clean up the free text like remove punctuations
2. Create train and test datasets
3. Feature for the models - jobtext (free text) and Label - Title
4. Fit training dataset to CountVectoriser
5. Run Gridsearch to get best Ngram paremeter
6. Run 12 classification models on train dataset and get accuracy scores to enable choose the best model
7. Train the model with the best accuracy score
8. Run the model with Test dataset
9. Check accuracy of the model

In [2]:

```

import requests
from bs4 import BeautifulSoup
import lxml
import numpy as np
import pandas as pd
from scipy.stats import norm
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
import re
# Getting that SKLearn Dataset
from sklearn.datasets import fetch_20newsgroups
%matplotlib inline

# machine learning
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC, LinearSVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import Perceptron
from sklearn.linear_model import SGDClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.naive_bayes import MultinomialNB

# Clustering
from sklearn.cluster import KMeans
from sklearn.preprocessing import normalize
from sklearn.metrics import silhouette_score

#Gridsearch and scoring
from sklearn.grid_search import GridSearchCV
from sklearn import metrics
from sklearn.model_selection import cross_val_score
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import accuracy_score

# Train test
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import HashingVectorizer, TfidfVectorizer
from sklearn.pipeline import make_pipeline
from sklearn.metrics import accuracy_score

```

C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\sklearn\cross\_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model\_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.

"This module will be removed in 0.20.", DeprecationWarning)

C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\sklearn\grid\_search.py:42: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model\_selection module into which all the refactored classes and functions are moved. This module will be removed in 0.20.

ions are moved. This module will be removed in 0.20.  
DeprecationWarning)

## Load data

In [3]:

```
dfe = pd.read_csv('./Data/jobsdf.csv', index_col=0)
```

In [4]:

```
dfe.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1768 entries, 0 to 1767
Data columns (total 15 columns):
category      1768 non-null object
company       1746 non-null object
jobclass      1768 non-null object
jobdate       1764 non-null object
jobid         1768 non-null int64
jobtext       1744 non-null object
location      1768 non-null object
subcategory   1768 non-null object
suburb        1296 non-null object
title         1768 non-null object
worktype      1764 non-null object
period        1768 non-null object
uppersal      289 non-null float64
lowersal      388 non-null float64
finalsal      388 non-null float64
dtypes: float64(3), int64(1), object(11)
memory usage: 221.0+ KB
```

In [5]:

```
# Classify salary
sal_type = []
for i in dfe.finalsal:
    if i > 0:
        sal_type.append(1)
    else:
        sal_type.append(0)
dfe['saltype'] = sal_type
```

In [6]:

```
# Cleanup data
import string
dfe['jobtext'].replace(np.nan, 'NA', inplace=True)
dfe['jobtext'].replace('\n', '', inplace=True)
dfe['jobtext'] = dfe['jobtext'].apply(lambda x: ''.join([i for i in x
                                                         if i not in string.punctuation]))
dfe['jobtext'].str.replace(r'^([a-z0-9_\-]+)@([\da-z\-\_]+\.[a-z\-\_]{2,6})$/', '')
```

Out[6]:

```
0      Offices in the USA UK and Australia Predictiv...
1      About SEEK SEEK is a diverse group of compan...
2      Associate Data Scientist Applied Machine Lear...
3      Data Scientist Applied Machine Learning About...
4      Ignite Data Solutions are a data and analytic...
5      Our client is a national reputable and wellkn...
6      We are seeking up to 2 Senior data scientists ...
7      Our client is an industry pioneering Customer ...
8      Do you want work in a place where Your supe...
9      Overview This multinational FinTech startup b...
10     Our client is one of Brisbanes fastest growing...
11     About the business and the role My client i...
12     Are you curious about how digital is changing...
13     Our client is one of Brisbanes leading Big Dat...
14     Offices in the USA UK and Australia Predictiv...
15     Are you a passionate Data Scientist looking t...
16     Data Scientist Our client an innovative sta...
17     Data Scientist We are currently recruiting...
18     My client is a marketleading IT ServicesSI who...
19     The organization is seeking a highly experien...
20     We are looking for a Data Scientist to join th...
21     We are a world leading Health Care organisat...
22     My client is a marketleading IT ServicesSI who...
23     Researcher Data Scientist AI Machine Learning...
24     About the Company One of Australia's leading ...
25     Data3 Limited DTL is an ASX listed company th...
26     Required skills A minimum of 5 years exper...
27     The Company My client is a wellestablished ...
28     A leading consulting company is driving a tra...
29     My client is a marketleading IT ServicesSI who...
...
1738    Whats it like working in the Engineering tea...
1739    FIRESOFT Consulting is a multiaward winning ...
1740    Our client is the largest and fastest growing...
1741    Melbourne PFM Head Office Supportive team en...
1742    Our Client is a global software vendor th...
1743    Looking to speak with the highest calibre of ...
1744    The role Our client a household Australian na...
1745    This is a precursor to improving the experien...
1746    Technology Solutions Manager My client op...
1747    This is a precursor to improving the experien...
1748    We are looking for an agile astute and affabl...
1749    My client a well known enterprise Retail orga...
1750    Global Company Market Leader Be part of a ...
1751    This is a precursor to improving the experien...
1752    Since 2008 Servian has been known to be one o...
1753    Servian was established in 2008 and has consi...
1754    Since 2008 Servian has been known to be one o...
1755    Represent one of the true cloudbased ERP sol...
```

```

1756 Progressive in partnership with one of the wo...
1757 Bluefin Resources are working with a prestigi...
1758 Based in Crows Nest acidgreen are an multiawa...
1759 Major project great team environment Permane...
1760 Excellent Corporate Package Super Company C...
1761 Are you a talented and skilled professional l...
1762 ABOUT USGAPbuster Worldwide is a global leade...
1763 We are currently recruiting for a Data Analys...
1764 Junior Sales Executive - Corporate Event Indu...
1765 The financial markets business of this global...
1766 Leading and managing an offshore and onshore ...
1767 South Eastern Sydney Local Health District SE...

```

Name: jobtext, Length: 1768, dtype: object

In [7]:

```

# Train and Test datasets FOR TITLES (SALARY TYPE)
X = dfe['jobtext']
#y=dfe.title
y=dfe.saltype
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.3, random_state=101)

```

In [8]:

```

# # Establish baseline
base_line = np.mean(dfe.finalsal)
base_line

```

Out[8]:

129809.76417525773

In [9]:

```

# Fit training dataset
vect= CountVectorizer(stop_words='english', ngram_range=(1,2))
vect.fit(X_train)

```

Out[9]:

```

CountVectorizer(analyzer='word', binary=False, decode_error='strict',
dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',
lowercase=True, max_df=1.0, max_features=None, min_df=1,
ngram_range=(1, 2), preprocessor=None, stop_words='english',
strip_accents=None, token_pattern='(?u)\\b\\w\\w+\\b',
tokenizer=None, vocabulary=None)

```

In [10]:

```

# Create DataFrame with vectorised words
XX_train = pd.DataFrame(vect.transform(X_train).todense(), columns= vect.get_feature_names(

```

In [11]:

```
# Check top 20 words
word_counts = XX_train.sum(axis=0)
word_counts.sort_values(ascending = False).head(20)
```

Out[11]:

data	7531
experience	4597
business	4384
skills	2377
role	2214
team	2132
work	1830
management	1639
working	1603
strong	1435
development	1389
analysis	1268
ability	1261
reporting	1259
information	1241
analytics	1228
apply	1202
support	1152
intelligence	1106
solutions	1058

dtype: int64

In [12]:

```
# Initialise NLTK word tokenizer
from nltk import word_tokenize
from nltk.stem import WordNetLemmatizer
class LemmaTokenizer(object):
    def __init__(self):
        self.wnl = WordNetLemmatizer()
    def __call__(self, articles):
        return [self.wnl.lemmatize(t) for t in word_tokenize(articles)]
```

**Run Gridsearch to get the best parameter for Ngram Range**

In [ ]:

```
# Run Gridsearch to get the best parameter for Ngram Range
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline
pipeline = Pipeline([
    ('vect', CountVectorizer()),
    ('mult', LogisticRegression()),
])
CV_parameters = { 'vect__ngram_range':((1, 1),(1, 2), (2, 2), (1, 3)),
                  'vect__stop_words': ['english'],
                }
CV_gridsearch = GridSearchCV(pipeline, CV_parameters, n_jobs=-1, verbose=1)

CV_gridsearch.fit(X_train, y_train)
print('Gridsearch Best Score: ', CV_gridsearch.best_score_, '\n')
print('Gridsearch Best Parameters: \n', CV_gridsearch.best_params_, '\n')
predictions = CV_gridsearch.predict(X_test)
print('Accuracy Score: ', accuracy_score(y_test, predictions), '\n')
print('Classification Report: \n', classification_report(predictions, y_test), '\n')
print('Confusion Matrix: \n', confusion_matrix(predictions, y_test), '\n')
```

Fitting 3 folds for each of 4 candidates, totalling 12 fits

```
C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\sklearn\model_selection
\_split.py:605: Warning: The least populated class in y has only 1 members,
which is too few. The minimum number of members in any class cannot be less
than n_splits=3.
% (min_groups, self.n_splits)), Warning)
```

In [23]:

```
# Check gridsearch output parameters
# print('Gridsearch Best Score: ', CV_gridsearch.best_score_, '\n')
# print('Gridsearch Best Parameters: \n', CV_gridsearch.best_params_, '\n')
# predictions = CV_gridsearch.predict(X_test)
# print('Accuracy Score: ', accuracy_score(y_test, predictions), '\n')
# print('Classification Report: \n', classification_report(predictions, y_test), '\n')
# print('Confusion Matrix: \n', confusion_matrix(predictions, y_test), '\n')
```

Out[23]:

```
<function __main__.run_Gridsearch>
```

**Run 12 classification models on train dataset and get accuracy scores to choose the best model**

In [122]:

```
#Creates a function to try the different models and print out the score
from sklearn.tree import DecisionTreeClassifier
def Models_Acc(XTrain, yTrain, XTest, yTest, ngramRange=(1, 2)):
    for i in range(1,12):
        if i == 1:
            model = make_pipeline(HashingVectorizer(stop_words='english',
                                                    non_negative=True,
                                                    n_features=2**16),
                                LogisticRegression())

            print ('-----')
            print ('Hashing Vectorizer and Logistic Regression')

        elif i == 2:
            model = make_pipeline(TfidfVectorizer(stop_words='english',
                                                  sublinear_tf=True,
                                                  max_df=0.5,
                                                  max_features=1000,
                                                  ngram_range=ngramRange),
                                LogisticRegression())

            print ('-----')
            print ('TfidfVectorizer using',format(ngramRange), 'ngram and Logistic Regression')
        elif i == 3:
            model = make_pipeline(CountVectorizer(stop_words='english', ngram_range=ngramRange),
                                LogisticRegression())

            print ('-----')
            print ('CountVectorizer with', format(ngramRange), 'ngram and Logistic Regression')

        elif i == 4:
            model = make_pipeline(CountVectorizer(stop_words='english', ngram_range=ngramRange),
                                LogisticRegression())

            print ('-----')
            print ('CountVectorizer with LemmaTokenizer and ', format(ngramRange), 'ngram and Logistic Regression')

        elif i == 5:
            model = make_pipeline(CountVectorizer(stop_words='english', ngram_range=ngramRange),
                                DecisionTreeClassifier(),

                                )

            print ('-----')
            print ('CountVectorizer with ', format(ngramRange), ' and MultinomialNB Regression')

        elif i == 6:
            model = make_pipeline(CountVectorizer(stop_words='english', ngram_range=ngramRange),
                                RandomForestClassifier(),

                                )

            print ('-----')
            print ('CountVectorizer with ', format(ngramRange), ' and Random Forest Regression')
        elif i == 7:
            model = make_pipeline(CountVectorizer(stop_words='english', ngram_range=ngramRange),
                                DecisionTreeClassifier())

            print ('-----')
            print ('CountVectorizer with ',format(ngramRange), 'and Decision Tree')

        elif i == 8:
```



```

model = make_pipeline(CountVectorizer(stop_words='english', ngram_range=ngramRange),
                      AdaBoostClassifier())

print ('-----')
print ('CountVectorizer with ',format(ngramRange), ' and AdaBoost')
elif i == 9:

    model = make_pipeline(TfidfVectorizer(stop_words='english',
                                          sublinear_tf=True,
                                          max_df=0.5,
                                          max_features=1000,
                                          ngram_range=ngramRange),
                          DecisionTreeClassifier())

    print ('-----')
    print ('TfidfVectorizer using ',format(ngramRange), ' ngram and DecisionTreeClassifier')

elif i == 10:
    model = make_pipeline(TfidfVectorizer(stop_words='english',
                                          sublinear_tf=True,
                                          max_df=0.5,
                                          max_features=1000,
                                          ngram_range=ngramRange),
                          RandomForestClassifier())

    print ('-----')
    print ('TfidfVectorizer using ',format(ngramRange), ' ngram and RandomForestClassifier')

elif i == 11:
    model = make_pipeline(CountVectorizer(stop_words='english', ngram_range=ngramRange),
                          SVC(kernel='linear'))

    print ('-----')
    print ('CountVectorizer with LemmaTokenizer with ',format(ngramRange), ' and SVM')

model.fit(XTrain, yTrain)
y_pred = model.predict(XTest)
print ('Accuracy:',accuracy_score(yTest, y_pred))

```

In [123]:

```
# Check accuracy of models
```

```
Models_Acc(X_train, y_train, X_test, y_test, ngramRange=(2, 2))
```

```
-----  
-----  
Hashing Vectorizer and Logistic Regression
```

```
C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\sklearn\feature_extraction\hashing.py:94: DeprecationWarning: the option non_negative=True has been deprecated in 0.19 and will be removed in version 0.21.  
  " in version 0.21.", DeprecationWarning)
```

```
C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\sklearn\feature_extraction\hashing.py:94: DeprecationWarning: the option non_negative=True has been deprecated in 0.19 and will be removed in version 0.21.  
  " in version 0.21.", DeprecationWarning)
```

```
C:\Users\Vinita Auplish\Anaconda3\lib\site-packages\sklearn\feature_extraction\hashing.py:94: DeprecationWarning: the option non_negative=True has been deprecated in 0.19 and will be removed in version 0.21.  
  " in version 0.21.", DeprecationWarning)
```

```
Accuracy: 0.7740112994350282  
-----  
-----
```

```
TfidfVectorizer using (2, 2) ngram and Logistic Regression
```

```
Accuracy: 0.7871939736346516  
-----  
-----
```

```
CountVectorizer with (2, 2) ngram and Logistic Regression
```

```
Accuracy: 0.839924670433145  
-----  
-----
```

```
CountVectorizer with LemmaTokenizer and (2, 2) ngram and Logistic Regression  
n
```

```
Accuracy: 0.8361581920903954  
-----  
-----
```

```
CountVectorizer with (2, 2) and MultinomialNB Regression
```

```
Accuracy: 0.8022598870056498  
-----  
-----
```

```
CountVectorizer with (2, 2) and Random Forest Regression
```

```
Accuracy: 0.8361581920903954  
-----  
-----
```

```
CountVectorizer with (2, 2) and Decision Tree
```

```
Accuracy: 0.7966101694915254  
-----  
-----
```

```
CountVectorizer with (2, 2) and AdaBoost
```

```
Accuracy: 0.775894538606403  
-----  
-----
```

```
TfidfVectorizer using (2, 2) ngram and DecisionTreeClassifier
```

```
Accuracy: 0.7532956685499058  
-----  
-----
```

```
TfidfVectorizer using (2, 2) ngram and RandomForestClassifier
```

```
Accuracy: 0.8173258003766478  
-----
```

-----

CountVectorizer with LemmaTokenizer with (2, 2) and SVM  
Accuracy: 0.8305084745762712

## Run Logistic Regression for Title (Salary Type) as that has given best accuracy score

In [13]:

```
# Lets use the stop_words argument to remove words like "and, the, a"
cvec = CountVectorizer(stop_words='english', ngram_range=(2,2))

# Fit our vectorizer using our train data
cvec.fit(X_train)
# Write vectorised words back to the dataframe
X_train = pd.DataFrame(cvec.transform(X_train).todense(),
                       columns=cvec.get_feature_names())
X_test = pd.DataFrame(cvec.transform(X_test).todense(),
                      columns=cvec.get_feature_names())

# RANDOM FOREST TAKES TOO LONG
# Import the model we are using
# from sklearn.ensemble import RandomForestRegressor
# # Instantiate model with 1000 decision trees
# rf = RandomForestRegressor(n_estimators = 1000, random_state = 42)
# rf.fit(X_train, y_train)
# rf.score(X_test, y_test)
# Import and fit our logistic regression and test it too
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr.fit(X_train, y_train)
lr.score(X_test, y_test)
```

Out[13]:

0.839924670433145

In [14]:

```
# Generate predictions from Logistic Regression
predictions = lr.predict(X_test)
print('Confusion Matrix: \n', confusion_matrix(y_test,predictions),'\n')
print('Classification Report: \n', classification_report(y_test,predictions),'\n')
```

Confusion Matrix:

```
[[399  11]
 [ 74  47]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.97	0.90	410
1	0.81	0.39	0.53	121
avg / total	0.84	0.84	0.82	531

In [15]:

```
# List top 30 predicted values from Logistic Regression
columns = np.array(cvec.get_feature_names())
def list_important_pred_var(data):
    important_pred_var = pd.DataFrame(data, columns = ["coef"], index = columns)
    return important_pred_var.sort_values(["coef"], ascending = False)[:30]

list_important_pred_var(lr.coef_.T)
```

Out[15]:

	coef
experience strong	0.365296
data engineer	0.336160
like apply	0.323484
confidence danieletegroupcomau	0.321953
danieletegroupcomau hot	0.321953
data modelling	0.315801
financial services	0.300634
send resume	0.299526
apply button	0.286014
initial month	0.284226
purpose position	0.275545
finance business	0.259355
development experience	0.257839
apply updated	0.245693
workforce planning	0.243020
sydney cbd	0.236132
analytics reporting	0.226177
based sydney	0.225985
work closely	0.225204
relevant experience	0.223625
internal stakeholders	0.222484
ability interpret	0.217470
ad hoc	0.214430
business partner	0.212060
apply feel	0.209412
data engineering	0.209228
high level	0.207831
data entry	0.203614
experience needed	0.203610
site business	0.196369

In [16]:

predictions

Out[16]:

```
array([[0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
        0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
        1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
        0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
        0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0,
        0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
        0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0,
        0, 0, 0], dtype=int64)
```

In [17]:

```
predictions.sum()
```

Out[17]:

58

In [ ]: