

Capstone Project - Telecom Network Service Disruption

Build Models with FULL dataset

This is Part 3 of the project documentation. (File Name 03-FINAL-TelstraModels-FullDataset)

Please see Part 1 for Project Details and Executive Summary. (File Name 01-FINAL -TelstraEDA)

Develop Models

Data

- Dependent variable is a 3 class variable **fault_severity**
- 5 features used **log_feature** (386 classes), **resource_type** (3 classes), **event_type** (6 classes), **location** (32 classes), **severity_type** (5 classes),

Estimators

- 5 classifier estimators have been selected for comparison - RandomForest Classifier, KNeighbors Classifier, DecisionTress Classifier, AdaBoost Classifier, GradientBoostClassifier

Hyperparameters

- *Hyperparameters for all estimators derived using `sklearn.model_selection.RandomizedSearchCV`*
Please see file 02-FINAL-TelstraHyperparameters for details on optimization with hyperparameters

Scenarios

- Each estimator was run for the following scenarios
- All features used - hyperparameters given
- All features used - default hyperparameters

Results

Estimator	All features			Features except logfeature	
	Default params	Optimum params		Default params	Optimum params
Random Forest Classifier	72.28%	75.49%		59.95%	62.53%
KNeighbors Classifier	71.33%	73.81%		59.37%	63.66%
DecisionTree Classifier	71.38%	70.11%		58.83%	58.19%
AdaBoost Classifier	72.42%	73.81%		66.23%	65.73%
GradientBoost Classifier	76.16%	76.48%		66.41%	65.96%
Baseline Accuracy	64.82%				

As can be seen, there is an improvement in accuracy, across all estimators except DecisionTree, when hyperparameters are provided to the model.

Top 20 features by Estimator

Random Forest		Decision Tree		AdaBoost		GradientBoost	
log_feature203	0.117975	log_feature203	0.196701	log_feature203	0.1	log_feature203	0.124773
log_feature82	0.083862	severity_type_1	0.058009	log_feature170	0.06	log_feature170	0.034011
log_feature170	0.03621	log_feature82	0.050477	resource_type_RT8	0.06	log_feature202	0.03206
log_feature54	0.033386	log_feature170	0.044608	event_type_OTH	0.04	log_feature209	0.024589
log_feature232	0.027159	log_feature54	0.02519	log_feature202	0.04	log_feature232	0.024231
log_feature312	0.022933	log_feature312	0.024538	location_995	0.02	log_feature312	0.023538
event_type_OTH	0.022118	log_feature80	0.022101	location_OTH	0.02	log_feature73	0.023496
log_feature80	0.021604	log_feature68	0.019273	event_type_ET11	0.02	log_feature82	0.018607
log_feature68	0.020152	log_feature232	0.017752	event_type_ET34	0.02	log_feature171	0.018412
log_feature71	0.018804	resource_type_OTH	0.015536	event_type_ET35	0.02	log_feature155	0.016335
location_OTH	0.016184	event_type_OTH	0.014884	severity_type_1	0.02	log_feature179	0.016276
event_type_ET15	0.016145	log_feature73	0.014877	log_feature193	0.02	severity_type_1	0.014595
event_type_ET34	0.015661	log_feature71	0.013795	log_feature195	0.02	log_feature134	0.01443
severity_type_1	0.015401	log_feature171	0.012685	log_feature196	0.02	log_feature315	0.014178
log_feature313	0.014771	log_feature315	0.012159	log_feature205	0.02	log_feature70	0.014025
log_feature201	0.014159	log_feature193	0.011945	log_feature140	0.02	log_feature368	0.013464
log_feature193	0.013333	log_feature201	0.011271	log_feature209	0.02	log_feature227	0.012689
severity_type_2	0.012206	log_feature291	0.011234	log_feature212	0.02	log_feature314	0.012604
log_feature73	0.011528	event_type_ET11	0.009957	log_feature319	0.02	log_feature54	0.012336
resource_type_RT8	0.011016	event_type_ET15	0.00971	log_feature295	0.02	event_type_OTH	0.012102

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sys
import time
```

Load Data

In [100]:

```
# Full dataset
train_id = pd.read_csv('./Data/train_id.csv')
test_id = pd.read_csv('./Data/test_id.csv')
```

In [101]:

```
print('Dataframe train - number of rows columns', train_id.shape)
print('Dataframe test - number of rows columns', test_id.shape)
```

Dataframe train - number of rows columns (7381, 434)

Dataframe test - number of rows columns (11171, 434)

Develop Models

In [102]:

```
#####
# Import classifier model modules
#####
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.svm import SVC
#Gridsearch and scoring
from sklearn.grid_search import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn import metrics
from sklearn.model_selection import cross_val_score
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
```

In [103]:

```
#####
# Prepare data for modelling
#####
# Create target dataframe
target = train_id['fault_severity']
# Drop fault_severity from train so we don't have to specify column names
train_id.drop(['fault_severity'], axis=1, inplace=True)
test_id.drop(['fault_severity'], axis=1, inplace=True)
train_id.drop(['id'], axis=1, inplace=True)
test_id.drop(['id'], axis=1, inplace=True)
```

In [104]:

```
#####
# Assign values to lists
#####
XTrain = train_id
yTrain = target
XTest = test_id
# Split training dataset into "train" and "validate" to compare model accuracy
X_train, X_test, y_train, y_test = train_test_split(XTrain, yTrain, test_size=0.3, random_st
```

In [105]:

```
#####
# Baseline accuracy on the percentage for largest class in training dataset
#####
baseline = target.value_counts()/len(target)
print('Baseline accuracy', round(baseline[0]*100,2))
```

Baseline accuracy 64.82

In [106]:

```
#####
# Run models on "train" and "test" and get relative scores
#####
# Random Forest Classifier with params
###
model = RandomForestClassifier(bootstrap=False,max_depth=70,max_features='auto',min_samples
                                n_estimators=1200)

model.fit(X_train, y_train)
y_pred = model.predict(X_test)
x = model.feature_importances_
print ('Random Forest Classifier accuracy - params:',round(accuracy_score(y_test, y_pred),4))
###
# Random Forest Classifier default
###
model = RandomForestClassifier()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print ('Random Forest Classifier accuracy - default:',round(accuracy_score(y_test, y_pred),4))
###
# Model feature importance
###
x = model.feature_importances_
zRF = pd.Series(data=x,index=X_train.columns)
print('Top 20 features')
print(zRF.sort_values(ascending=False).head(20))
```

```
Random Forest Classifier accuracy - params: 0.7549
Random Forest Classifier accuracy - default: 0.7228
Top 20 features
log_feature203      0.117975
log_feature82       0.083862
log_feature170      0.036210
log_feature54       0.033386
log_feature232      0.027159
log_feature312      0.022933
event_type_OTH      0.022118
log_feature80       0.021604
log_feature68       0.020152
log_feature71       0.018804
location_OTH        0.016184
event_type_ET15     0.016145
event_type_ET34     0.015661
severity_type_1     0.015401
log_feature313      0.014771
log_feature201      0.014159
log_feature193      0.013333
severity_type_2     0.012206
log_feature73       0.011528
resource_type_RT8   0.011016
dtype: float64
```

In [107]:

```
###  
# KNeighbors Classifier params  
###  
model = KNeighborsClassifier(algorithm='kd_tree',metric='minkowski',leaf_size=10,p=4,weight  
model.fit(X_train, y_train)  
y_pred = model.predict(X_test)  
print ('KNeighbors Classifier accuracy - params:',round(accuracy_score(y_test, y_pred),4))  
###  
# KNeighbors Classifier params default  
###  
model = KNeighborsClassifier()  
model.fit(X_train, y_train)  
y_pred = model.predict(X_test)  
print ('KNeighbors Classifier accuracy - default:',round(accuracy_score(y_test, y_pred),4))  
###  
# Model feature importance not available in KNN  
###
```

KNeighbors Classifier accuracy - params: 0.7381

KNeighbors Classifier accuracy - default: 0.7133

In [109]:

```

###
# DecisionTree Classifier params
###
model = DecisionTreeClassifier(class_weight='balanced',criterion='gini',max_features=None,s
model.fit(X_train,y_train)
y_pred = model.predict(X_test)
print ('DecisionTree Classifier accuracy - params:',round(accuracy_score(y_test, y_pred),4))
###
# DecisionTree Classifier default
###
model = DecisionTreeClassifier()
model.fit(X_train,y_train)
y_pred = model.predict(X_test)
print ('DecisionTree Classifier accuracy - default:',round(accuracy_score(y_test, y_pred),4))
###
# Model feature importance
###
x = model.feature_importances_
zDT = pd.Series(data=x,index=X_train.columns)
print('Top 20 features')
print(zDT.sort_values(ascending=False).head(20))

```

DecisionTree Classifier accuracy - params: 0.7011

DecisionTree Classifier accuracy - default: 0.7138

Top 20 features

log_feature203	0.196701
severity_type_1	0.058009
log_feature82	0.050477
log_feature170	0.044608
log_feature54	0.025190
log_feature312	0.024538
log_feature80	0.022101
log_feature68	0.019273
log_feature232	0.017752
resource_type_OTH	0.015536
event_type_OTH	0.014884
log_feature73	0.014877
log_feature71	0.013795
log_feature171	0.012685
log_feature315	0.012159
log_feature193	0.011945
log_feature201	0.011271
log_feature291	0.011234
event_type_ET11	0.009957
event_type_ET15	0.009710

dtype: float64

In [110]:

```

####
# AdaBoost Classifier params
####
model = AdaBoostClassifier(algorithm='SAMME.R',learning_rate=1.0,n_estimators=90,random_state=0)
model.fit(X_train,y_train)
y_pred = model.predict(X_test)
print ('AdaBoost Classifier accuracy - params:',round(accuracy_score(y_test, y_pred),4))
####
# AdaBoost Classifier default
####
model = AdaBoostClassifier()
model.fit(X_train,y_train)
y_pred = model.predict(X_test)
print ('AdaBoost Classifier accuracy - default:',round(accuracy_score(y_test, y_pred),4))
####
# Model feature importance
####
x = model.feature_importances_
zAB = pd.Series(data=x,index=X_train.columns)
print('Top 20 features')
print(zAB.sort_values(ascending=False).head(20))

```

AdaBoost Classifier accuracy - params: 0.7381
 AdaBoost Classifier accuracy - default: 0.7242

Top 20 features

log_feature203	0.10
log_feature170	0.06
resource_type_RT8	0.06
event_type_OTH	0.04
log_feature202	0.04
location_995	0.02
location_OTH	0.02
event_type_ET11	0.02
event_type_ET34	0.02
event_type_ET35	0.02
severity_type_1	0.02
log_feature193	0.02
log_feature195	0.02
log_feature196	0.02
log_feature205	0.02
log_feature140	0.02
log_feature141	0.02

In [111]:

```

###
# GradientBoost Classifier params
###
model = GradientBoostingClassifier(criterion='friedman_mse',init=None,learning_rate=0.1,max
                                   n_estimators=150,random_state=88)
model.fit(X_train,y_train)
y_pred = model.predict(X_test)
print ('GradientBoost Classifier accuracy - params:',round(accuracy_score(y_test, y_pred),4))
###
# GradientBoost Classifier params
###
model = GradientBoostingClassifier()
model.fit(X_train,y_train)
y_pred = model.predict(X_test)
print ('GradientBoost Classifier accuracy - default:',round(accuracy_score(y_test, y_pred),4))
###
# Model feature importance
###
x = model.feature_importances_
zGB = pd.Series(data=x,index=X_train.columns)
print('Top 20 features')
print(zGB.sort_values(ascending=False).head(20))

```

GradientBoost Classifier accuracy - params: 0.7616

GradientBoost Classifier accuracy - default: 0.7648

Top 20 features

log_feature203	0.124773
log_feature170	0.034011
log_feature202	0.032060
log_feature209	0.024589
log_feature232	0.024231
log_feature312	0.023538
log_feature73	0.023496
log_feature82	0.018607
log_feature171	0.018412
log_feature155	0.016335
log_feature179	0.016276
severity_type_1	0.014595
log_feature134	0.014430
log_feature315	0.014178
log_feature70	0.014025
log_feature368	0.013464
log_feature227	0.012689
log_feature314	0.012604
log_feature54	0.012336
event_type_OTH	0.012102

dtype: float64

In [115]:

```
X_test['log_feature203'][X_test.log_feature203 > 0].count()
```

Out[115]:

337

In [116]:

```
X_test['log_feature170'][X_test.log_feature170 > 0].count()
```

Out[116]:

177

In [117]:

```
X_test['log_feature82'][X_test.log_feature82 > 0].count()
```

Out[117]:

410

In []: